

IBM University Engagement Project Submission

“Exploring the Power of Agentic AI with IBM Granite and LangFlow”

MindGuard AI – Agentic AI for Mental Health Awareness & Suicide Prevention

Domain of Project – Healthcare – Mental Health & Suicide Prevention

Student Name	Recent Passport Photo	Email ID	Mobile number [WhatsApp]
Saksham Gupta		sakshamgupta12312030@gmail.com	9389770509

Problem Statement -

Problem Statement : Agentic AI for Mental Health Awareness & Suicide Prevention

The Challenge

Mental health issues such as **stress, anxiety, depression, loneliness, and suicidal thoughts** are increasing worldwide. Many people hesitate to seek professional help due to stigma, lack of awareness, limited access to mental health resources, and delayed intervention. Existing solutions are often reactive and fail to provide **continuous emotional support, early risk detection, personalized self-care guidance, and secure journaling**.

The Objective

Build an **AI-powered Agentic Mental Health Assistant (MindGuard AI)** that provides **empathetic conversations, emotional wellness support, early distress detection, and personalized self-care guidance** for users. The solution should include:

- **AI Mental Health Chatbot** – Provide empathetic conversations using IBM watsonx.ai while maintaining user safety.
- **Mood Tracking & Analytics** – Allow users to record daily moods and visualize emotional trends over time.
- **Personal Journal** – Enable secure journaling with Cloudant storage for emotional reflection and progress tracking.
- **Self-Care Exercises** – Recommend breathing exercises, meditation, stretching, gratitude practice, and mindfulness sessions.
- **Mental Health Resource Center (RAG)** – Retrieve trusted mental health information and wellness guidance using Retrieval-Augmented Generation.
- **Emergency SOS Support** – Provide immediate access to emergency contacts, crisis helplines, and safety guidance during high-risk situations.
- **Safety Classification Agent** – Detect potentially harmful or suicidal conversations and provide safe, responsible responses with escalation guidance.
- **Visualization Dashboard** – Display mood history, journaling activity, wellness progress, and self-care engagement through interactive analytics.

Proposed Solution -

Proposed Solution: MindGuard AI – Agentic Mental Health Assistant

MindGuard AI is an **AI-powered mental health assistant** built using **IBM watsonx.ai, IBM Granite Models, IBM Cloudant, RAG, React, and Node.js**. It provides **empathetic conversations, mood tracking, secure journaling, personalized self-care guidance, and emergency support**.

At its core, the solution uses a **Retrieval-Augmented Generation (RAG)** pipeline to retrieve trusted mental health resources and generate **safe, context-aware, and empathetic responses** using IBM Granite models.

The system consists of the following key modules:

AI Mental Health Chatbot – Provides empathetic conversations and emotional support.

Mental Health RAG Agent – Retrieves trusted mental health knowledge and coping strategies.

Mood Tracking Agent – Records moods and visualizes emotional trends.

Personal Journal Agent – Stores journals securely using IBM Cloudant.

Self-Care Agent – Recommends breathing, meditation, mindfulness, and wellness exercises.

Safety Classification Agent – Detects high-risk conversations and generates responsible responses.

Emergency SOS Agent – Provides quick access to crisis helplines and emergency contacts.

The solution includes a **user-friendly dashboard** that visualizes mood history, journaling activity, and overall wellness progress.

IBM watsonx.ai powers conversational intelligence, **IBM Granite Models** handle reasoning, **IBM Cloudant** stores user data securely, and **RAG** enhances responses with trusted mental health knowledge.

Overall, this solution delivers an **intelligent, secure, and privacy-focused mental health assistant** that promotes **mental health awareness, early risk detection, emotional well-being, and timely support**.

Technology Used & Implementation

- **IBM Granite Model** – Used for natural language understanding, empathetic conversations, reasoning, and response generation.
- **IBM watsonx.ai** – Powers AI model inference, prompt execution, and conversational intelligence.
- **IBM Cloud** – Provides a secure and scalable cloud environment for deploying the application and AI services.
- **IBM Cloudant** – NoSQL database used to securely store user journals, mood history, conversations, and application data.
- **Retrieval-Augmented Generation (RAG)** – Retrieves trusted mental health resources to generate accurate and context-aware responses.
- **React.js + Node.js + Express.js** – Used to build the interactive frontend and backend APIs for the application.

IBM Granite Model

IBM watsonx

Upgrade ⓘ 🔔

Saksham Gupta's Account ▾

Sydney ▾

96

Resource hub ⓘ

Foundation models ▾ Samples

▼ 🔍 Search resources

Manage models ⓘ

Filter

Task

☐ Classification (74)
☐ Code explanation (18)
☐ Code fixing (18)
☐ Code generation (18)
☐ Code generation and conversion (67)
[View all \(12\)](#)

Provider

☐ BigScience (1)
☐ Code Llama (1)
☐ Google (1)
☐ IBM (27)
☐ LumioOpen (1)
[View all \(12\)](#)

Function & modality

☐ Audio (1)
☐ Image (20)
☐ Language (60)
☐ Video (1)

Model type

☒ All (81)
☐ Deploy on demand (76)
☐ Provided model (6)

llama-3-3-70b-instruct

🟢 Ready to use

The Meta Llama 3.3 multilingual large language model (LLM) is a pretrained and...

Provider: Meta Type: Provided model

language

mt0-xxl-13b

🟡 Not yet deployed

mt0-xxl (13B) is an instruction-tuned iteration on mt5. Like BLOOMZ, it was fin...

Provider: BigScience Type: Deploy on demand

language

codellama-34b-instruct-hf

🟡 Not yet deployed

Code Llama is a pretrained and fine-tuned generative text models with 34 billion...

Provider: Code Llama Type: Deploy on demand

language

deepseek-r1-distill-llama-70b

🟡 Not yet deployed

DeepSeek-R1-Distill models are fine-tuned based on open-source models, using...

Provider: deepseek-ai Type: Deploy on demand

language

deepseek-r1-distill-llama-8b

🟡 Not yet deployed

DeepSeek-R1-Distill models are fine-tuned based on open-source models, using...

Provider: deepseek-ai Type: Deploy on demand

language

gemma-4-31b-it

🟡 Not yet deployed

Gemma is a family of open models built by Google DeepMind. Gemma 4 models are...

Provider: Google Type: Deploy on demand

image

granite-13b-chat-v2

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-20b-code-base-sql-gen

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-20b-code-instruct

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-20b-multilingual

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-3-1-8b-base

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-3-1-8b-instruct

🟡 Not yet deployed

Granite-3.1-8B-Instruct is a 8B parameter long-context instruct model finetuned fro...

Provider: IBM Type: Deploy on demand

language

granite-3-2-8b-instruct

🟡 Not yet deployed

Granite-3.2 8B Instruct is a text-only model capable of reasoning which can be enable...

Provider: IBM Type: Deploy on demand

language

granite-3-3-2b-instruct

🟡 Not yet deployed

Granite-3.3-2B-Instruct is a 2B parameter long-context instruct model finetuned fro...

Provider: IBM Type: Deploy on demand

language

granite-3-3-8b-instruct

🟡 Not yet deployed

Granite-3.3-8B-Instruct is a 8B parameter long-context instruct model finetuned fro...

Provider: IBM Type: Deploy on demand

language

granite-3-8b-base

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-34b-code-instruct

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-3b-code-instruct

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-4-1-30b

🟡 Not yet deployed

Granite-4.1-30B is a 30B parameter long-context instruct model finetuned from...

Provider: IBM Type: Deploy on demand

language

granite-4-1-3b

🟡 Not yet deployed

Granite-4.1-3B is a 3B parameter long-context instruct model finetuned from...

Provider: IBM Type: Deploy on demand

language

granite-4-1-8b

🟡 Not yet deployed

Granite-4.1-8B is a 8B parameter long-context instruct model finetuned from...

Provider: IBM Type: Deploy on demand

language

granite-4-h-micro

🟡 Not yet deployed

Granite-4.0-H-Micro is a 3B parameter long-context instruct model finetuned fro...

Provider: IBM Type: Deploy on demand

language

granite-4-h-small

🟡 Not yet deployed

Granite-4.0-H-Small is a 30B parameter long-context instruct model finetuned fro...

Provider: IBM Type: Deploy on demand

language

granite-4-h-tiny

🟡 Not yet deployed

Granite-4.0-H-Tiny is a 7B parameter long-context instruct model finetuned fro...

Provider: IBM Type: Deploy on demand

language

granite-7b-lab

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-8b-code-instruct

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-8b-japanese

🟡 Not yet deployed

Granite models are designed to be used for a wide range of generative and non...

Provider: IBM Type: Deploy on demand

language

granite-speech-4-1-2b

🟡 Not yet deployed

Granite Speech 4.1 2B is a compact and efficient speech-language model...

Provider: IBM Type: Deploy on demand

language

au-syd.dai.cloud.ibm.com/wx/samples/.../deepseek-r1-distill-llama-70b-cura...

Desktop 1

IBM Cloudant

IBM Cloud

Search resources and products...

Q

Catalog

Manage

Saksham Gupta's Account

?

Resource list /

mindguard-cloudant ✓ [Add tags](#) [✎](#)

Details

Actions

Manage

Service credentials

Plan

Overview

Capacity

Docs

Launch Dashboard

Deployment details

CRN

crn:v1:bluemix:public:cloudantnosqldb:au-syd:a/d925754752db49faa89d6db78762e6a6:82ef3fca-ac30-4700-8c6e-da6814c13f1f::

Location

Sydney

External endpoint

<https://8f30ef42-e865-49a8-8357-fa76badb3e66-bluemix.cloudant.com>

External endpoint (preferred)

<https://8f30ef42-e865-49a8-8357-fa76badb3e66-bluemix.cloudantnosqldb.appdomain.cloud>

Authentication methods

[IBM Cloud IAM](#)

Activity Tracker event types

Management

Save

Disk encryption

Yes. Automatically generated disk encryption key.

Capacity details

Resource List

IBM Cloud

Search resources and products...

Catalog

Manage

Saksham Gupta's Account

Resource list

Create resource +

Name	Group	Location	Product	Status	Tags
m	Filter by group...	Filter...	Filter...	Filter...	Filter...
AI / Machine Learning (1 / 1)					
mindguard-runtime	Default	Sydney (au-syd)	watsonx.ai Runtime	Active	cpdaas
Databases (1 / 1)					
mindguard-cloudant	Default	Sydney (au-syd)	Cloudant	Active	-

20°

Windows Taskbar Icons

Search

ENG IN

3:19 PM
7/18/2026

IBM Bob

The screenshot displays the IBM Bob development environment, which is a web-based IDE. The interface is divided into several sections:

- EXPLORER (Left Sidebar):** Shows the project structure. The 'services' folder is expanded, showing files like `ragService.js`, `safetyClassifier.js`, and `watsonxCli.js`. Other folders include 'backend', 'src', 'tests', 'docs', 'frontend', 'infra', 'code-engine', and 'deploy.sh'.
- Editor (Center):** The main workspace shows the `ragService.js` file. The content is a Windows command prompt window with the following text:

```
Microsoft Windows [Version 10.0.26200.8655]
(c) Microsoft Corporation. All rights reserved.

E:\edunet\menu>
```
- TERMINAL (Bottom):** Displays the output of the application. It shows the RAG (Retrieval-Augmented Generation) process, including indexing knowledge files and running the backend. The output includes:

```
[RAG] Indexed "self-care-wellness.txt" - 15 chunks
[RAG] Indexed "suicide-prevention.txt" - 14 chunks
[RAG] Ready. Total indexed chunks: 48
[MindGuard] Backend running on http://localhost:3001
[MindGuard] Environment: development
```
- Right Sidebar:** Contains a 'Pasted text #1 - 51 lines' section with a token count of 234.2k / 270.0k and a cost of 21.52. Below this is a 'All tasks completed!' message with 14/14 tasks. A list of tasks is shown, including 'retrieve()', 'scoreChunk()', 'tokenise()', and 'ragAnswer()'. The 'Startup log you will now see' section shows the RAG initialization process, including indexing knowledge files and running the backend.

Project Dashboard

 **MindGuardAI**

 Chat

 Journal

 Mood Tracker

 Exercises

 Resources

 Emergency

IBM watsonx.ai Powered

You're never alone.


MindGuard AI provides emotional support, mental wellness guidance, trusted resources and early intervention.

[Start Conversation →](#)[Take Mental Health Check](#)


"Every small step matters."

Data is stored safely in IBM Cloudant.

 [SOS Help](#)




Project Chat Interface



AI Components Used-

- 1) **Chat Input** – Accepts user messages, mood updates, and journal entries.
- 2) **IBM watsonx.ai Agent** – Processes user input and generates safe, empathetic responses.
- 3) **IBM Granite Model – Granite-3-3-8B-Instruct** for reasoning, conversation, and mental health guidance.
- 4) **Chat Output** – Displays AI-generated responses, self-care suggestions, and wellness guidance.
- 5) **RAG Component** – Retrieves trusted mental health resources and coping strategies before generating responses.
- 6) **Cloudant Storage Component** – Securely stores user journals, mood history, conversations, and application data.

System Workflow-



Architecture blueprint

MindGuard AI

Agentic Mental Health Awareness & Suicide Prevention Assistant — Enterprise Architecture Blueprint



Role of Agentic AI in the solution

Role of Agentic AI in MindGuard AI


Agentic AI enables **MindGuard AI** to function as an intelligent, proactive mental health assistant rather than a simple chatbot. Using **IBM watsonx.ai** and **IBM Granite models**, it coordinates multiple specialized AI agents to provide empathetic conversations, emotional support, and personalized self-care guidance.

Each agent performs a specific role—such as **safety classification, mood analysis, RAG-based knowledge retrieval, journaling support, and self-care recommendations**—while working collaboratively. This **multi-agent architecture** ensures accurate, context-aware, and reliable responses.

Agentic AI also provides **context awareness**, understanding users' mood history, journal entries, and conversation context to generate personalized recommendations. It supports **real-time adaptation**, continuously adjusting responses based on the user's emotional state and interactions.

Additionally, the agents communicate with each other to improve decision-making, detect high-risk situations, and provide timely crisis support when required.


Overall, **Agentic AI** makes **MindGuard AI** proactive, empathetic, personalized, and safety-focused, helping users improve their mental well-being through intelligent assistance and early intervention.




Token Usage

12,003 token are used

Before Prompt



After the Prompt



Novelty and Uniqueness

Key Features / Unique Selling Points (USP) of MindGuard AI

The proposed **MindGuard AI** stands out by combining **Agentic AI**, **RAG**, and **personalized mental health support** into one intelligent system.

Multi-Agent Intelligence – Multiple AI agents collaborate to provide safe, empathetic, and context-aware mental health assistance.

RAG-Based Knowledge Retrieval – Retrieves trusted mental health information to reduce misinformation and improve response quality.

Personalized Emotional Support – Delivers recommendations based on mood history, journal entries, and conversation context.

Real-Time Risk Detection – Continuously analyzes user interactions to detect emotional distress or high-risk situations.

Integrated Wellness Features – Combines AI Chat, Mood Tracking, Journaling, Self-Care Exercises, and Emergency SOS in one platform.

Safety-First Design – Uses AI safety classification to provide responsible responses and crisis guidance.

Scalable & Enterprise-Ready – Built using **IBM watsonx.ai**, **Granite Models**, **Cloudant**, and **RAG** for secure and reliable deployment.

This makes MindGuard AI a proactive, empathetic, and intelligent mental health companion rather than just a simple chatbot.

Git Hub Link

<https://github.com/codecrafter34/MindGuard.git>

.

.

.

Future Scope

1. Wearable & Health Device Integration –

Integrate with smartwatches and fitness bands to monitor heart rate, sleep, stress levels, and activity. This will enable real-time mental wellness insights and personalized self-care recommendations.

2. Voice Assistant & Multilingual Support –

Add voice-based interaction and support multiple regional languages to make the platform more accessible and user-friendly for people from diverse backgrounds.

3. Professional Counselor Integration –

Enable secure connections with licensed psychologists, therapists, and emergency helplines for timely intervention and continuous mental health support.

4. Mobile App & Smart Notifications –

Develop Android/iOS applications with personalized reminders for mood tracking, journaling, breathing exercises, and medication or therapy schedules.

.

Certificates Earned:

Please add your credly certificate

In recognition of the commitment to achieve
professional excellence



Saksham Gupta

Has successfully satisfied the requirements for:

Getting Started with Artificial Intelligence



Issued on: Jun 18, 2026
Issued by: IBM SkillsBuild

Verify: <https://www.credly.com/badges/8e9c9ae5-9688-4782-94a0-0b13a6d10059>



IBM SkillsBuild

Completion Certificate



This certificate is presented to

Saksham Gupta

for the completion of

Exploring Quantum Computing

(ALM-COURSE_3825253)

According to the Adobe Learning Manager system of record

Completion date: 23 Jun 2026 (GMT)

Learning hours: 50 mins

IBM **SkillsBuild**

Completion Certificate



This certificate is presented to

Saksham Gupta

for the completion of

Journey to Cloud: Envisioning Your Solution

(ALM-COURSE_4058877)

According to the Adobe Learning Manager system of record

Completion date: 23 Jun 2026 (GMT)

Learning hours: 5 hrs

In recognition of the commitment to achieve
professional excellence



Saksham Gupta

Has successfully satisfied the requirements for:

Getting Started with Cybersecurity



Issued on: Jun 26, 2026
Issued by: IBM SkillsBuild

Verify: <https://www.credly.com/badges/d0a3e70c-758d-42e9-9261-3d10cec7131a>



IBM **SkillsBuild**

Completion Certificate



This certificate is presented to

Saksham Gupta

for the completion of

Level Up Cybersecurity with Generative AI

(ALM-COURSE_3947352)

According to the Adobe Learning Manager system of record

Completion date: 26 Jun 2026 (GMT)

Learning hours: 1 hr 30 mins

Certificates Earned:

IBM **SkillsBuild**

Completion Certificate



This certificate is presented to

Saksham

for the completion of

Lab: Troubleshoot Your Code Using IBM Bob

(ALM-COURSE_4071307)

According to the Adobe Learning Manager system of record

Completion date: 19 Jun 2026 (GMT)

Learning hours: 30 mins

Thank You!

Thank you for your time and interest.